

Lecture 10

R and the tidyverse // randomization

Ivan Rudik
AEM 6510

Roadmap

- What is R?
- What is the tidyverse?
- How do we import and manipulate data?

Our goal is to take a hands on approach to learning how we do environmental economics research

A good chunk of this lecture comes from Grant McDermott's [data science for economists](#) notes, and [RStudio education](#)

RStudio Cloud

Getting started

We will be using [rstudio.cloud](#) for our coding

Getting started

We will be using **rstudio.cloud** for our coding

Why?

Getting started

We will be using [rstudio.cloud](#) for our coding

Why?

You don't need to download/install anything

Getting started

We will be using [rstudio.cloud](#) for our coding

Why?

You don't need to download/install anything

I can prepare the packages and code and make it easy to download

Getting started

We will be using [rstudio.cloud](#) for our coding

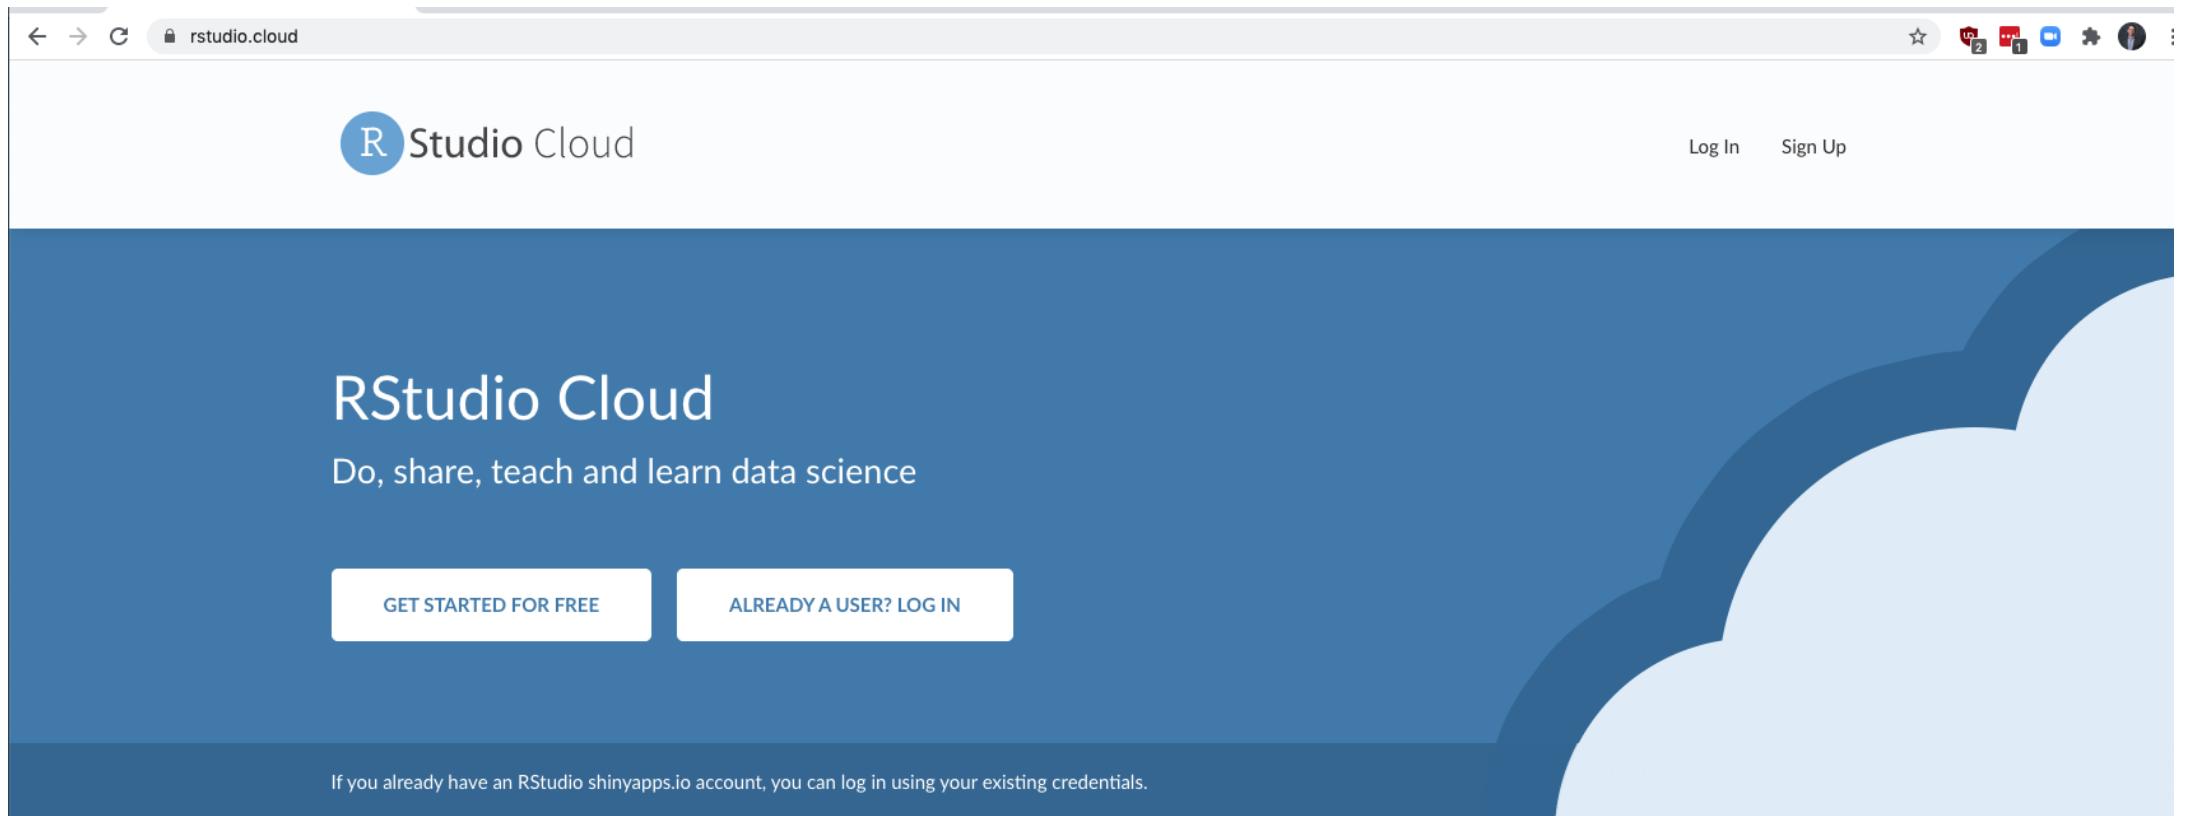
Why?

You don't need to download/install anything

I can prepare the packages and code and make it easy to download

Let's get everything going...

Getting started: login



The screenshot shows the RStudio Cloud homepage. At the top, there's a navigation bar with icons for back, forward, refresh, and search, followed by the URL 'rstudio.cloud'. On the right side of the bar are various browser extension icons. Below the bar, the RStudio Cloud logo is on the left, and 'Log In' and 'Sign Up' buttons are on the right. The main background is a large blue graphic featuring a white cloud shape on the right. The text 'RStudio Cloud' is prominently displayed in white, with the subtitle 'Do, share, teach and learn data science' below it. Two buttons are visible: 'GET STARTED FOR FREE' and 'ALREADY A USER? LOG IN'. At the bottom of the main section, a note says: 'If you already have an RStudio shinyapps.io account, you can log in using your existing credentials.'

Data science without the hardware hassles

RStudio Cloud is a lightweight, cloud-based solution that allows anyone to do, share, teach and learn data science online.

- Analyze your data using the RStudio IDE, directly from your browser.

\$ AVAILABLE PRICING PLANS

🕒 RSTUDIO CLOUD GUIDE

🌐 RSTUDIO.COM

Getting started: new RStudio project

The screenshot shows the RStudio Cloud interface. At the top, there's a navigation bar with 'F21 AEM 6510' (with a user icon), 'Projects' (which is selected and highlighted in blue), 'Members', and 'About'. On the far right of the top bar are icons for settings, a trash can, three dots for more options, and a user profile for 'Ivan Rudik'. Below the top bar, the main area is titled 'All Projects'. It includes filters for 'List' (set to 'All projects') and 'Sort' (set to 'By name'). A single project is listed: 'Untitled Project' by 'Ivan Rudik' (RStudio Project, created Oct 26, 2021 2:15 PM). To the right of the project list is a 'New Project' dropdown menu with three options: 'New RStudio Project' (selected, indicated by a blue background), 'New Jupyter Project', and 'New Project from Git Repository'.

F21 AEM 6510

Projects Members About

Ivan Rudik

All Projects

New Project

New RStudio Project

New Jupyter Project

New Project from Git Repository

Untitled Project

Ivan Rudik RStudio Project Created Oct 26, 2021 2:15 PM

Getting started: wait for deployment

The screenshot shows the RStudio Cloud interface with the URL `rstudio.cloud/project/1795327` in the address bar. The top navigation bar includes icons for back, forward, refresh, and user authentication. The main header says "Your Workspace / aem6510". On the left, a sidebar menu has "Your Workspace" selected. The menu items include "Spaces", "Learn" (with "Guide", "What's New", "Primers", and "Cheat Sheets"), "Help" (with "Current System Status" and "RStudio Community"), and "Info" (with "Plans & Pricing" and "Terms and Conditions"). In the center, a large circular progress bar indicates a deployment process, with the text "Deploying Project" and a blue progress bar at the bottom.

Click on class-code.Rproj

The screenshot shows the RStudio Cloud interface at rstudio.cloud/project/1795327. The workspace is named "aem6510".

Left Sidebar:

- Spaces: Your Workspace (selected), New Space.
- Learn: Guide, What's New, Primers, Cheat Sheets.
- Help: Current System Status, RStudio Community.
- Info: Plans & Pricing, Terms and Conditions.

Top Bar:

- File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, Help.
- Go to file/function, Addins.
- R 4.0.3.

Console Tab:

Restarting R session...

```
>  
> |
```

Environment Tab:

Import Dataset, Global Environment. Environment is empty.

Files Tab:

Name	Size	Modified
..		
.Rprofile	26 B	Oct 20, 2020, 1:48 PM
class-code.Rproj	205 B	Oct 20, 2020, 1:48 PM
code-here.R	101 B	Oct 20, 2020, 1:48 PM
renv		
renv.lock	17.6 KB	Oct 20, 2020, 1:48 PM

Click yes

The screenshot shows the RStudio Cloud interface. On the left, there's a sidebar with links for Spaces, Learn, Help, and Info. The main area has tabs for Console, Terminal, and Jobs. The Console tab shows the message "Restarting R session...". To the right, there are tabs for Environment, History, Connections, Git, and Tutorial. The Environment tab displays the message "Environment is empty". Below the tabs, there's a file browser with a "Confirm Open Project" dialog box. The dialog asks "Do you want to open the project /cloud/project/class-code?". It has "Yes" and "No" buttons. The file browser lists files: ".Rproj" (26 B, Oct 20, 2020, 1:48 PM), ".R" (205 B, Oct 20, 2020, 1:48 PM), ".R" (101 B, Oct 20, 2020, 1:48 PM), and "renv.lock" (17.6 KB, Oct 20, 2020, 1:48 PM). The top navigation bar shows the URL "rstudio.cloud/project/1795327" and the user "Ivan Rudik".

Quick intro to R

Arithmetic operations

R can do all the standard arithmetic operations

```
1+2 ## add
```

```
## [1] 3
```

```
6-7 ## subtract
```

```
## [1] -1
```

```
5/2 ## divide
```

```
## [1] 2.5
```

Logical operations

You also have logical operations

```
1 > 2
```

```
## [1] FALSE
```

```
(1 > 2) | (1 > 0.5) # / is the or operator
```

```
## [1] TRUE
```

```
(1 > 2) & (1 > 0.5) # & is the and operator
```

```
## [1] FALSE
```

Logical operations

We can negate expressions with: !

This is helpful for filtering data

```
is.na(1:10)
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
!is.na(1:10)
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

Logical operators

For value matching we use: `%in%`

To see whether an object is contained within (i.e. matches one of) a list of items, use `%in%`.

```
4 %in% 1:10
```

```
## [1] TRUE
```

```
4 %in% 5:10
```

```
## [1] FALSE
```

This is kind of like an `any` command in other languages

Logical operators

To evaluate whether two expressions are equal, we need to use **two** equal signs

```
1 = 1 ## This doesn't work
```

```
## Error in 1 = 1: invalid (do_set) left-hand side to assignment
```

```
1 == 1 ## This does.
```

```
## [1] TRUE
```

```
1 != 2 ## Note the single equal sign when combined with a negation.
```

```
## [1] TRUE
```

Assignment

In R, we can use either `=` or `←` to handle assignment.¹

¹ The `←` is really a `<` followed by a `-`. It just looks like an arrow because of the font on the slides.

Assignment

In R, we can use either `=` or `←` to handle assignment.¹

You can think of it as a (left-facing) arrow saying **assign in this direction**

¹ The `←` is really a `<` followed by a `-`. It just looks like an arrow because of the font on the slides.

Assignment

```
a ← 10 + 5  
a
```

```
## [1] 15
```

Assignment

You can also use `=` for assignment

```
b = 10 + 10  
b
```

```
## [1] 20
```

Which assignment operator to use?

Most R folks prefer `←` for assignment

In RStudio you can assign `←` to a hotkey to make using it as easy as `=`

Which assignment operator to use?

Most R folks prefer `←` for assignment

`=` has a specific role for evaluation *within* functions too

In RStudio you can assign `←` to a hotkey to make using it as easy as `=`

Which assignment operator to use?

Most R folks prefer `←` for assignment

`=` has a specific role for evaluation *within* functions too

It doesn't really matter though, other languages use `=` for both

In RStudio you can assign `←` to a hotkey to make using it as easy as `=`

Which assignment operator to use?

Most R folks prefer `←` for assignment

`=` has a specific role for evaluation *within* functions too

It doesn't really matter though, other languages use `=` for both

Use whatever you prefer, just be consistent

In RStudio you can assign `←` to a hotkey to make using it as easy as `=`

Help

If you are struggling with a (named) function or object in R, simply type ?
commandhere

?Negate

Help

Also try `vignette()` for a more detailed introduction to many packages

```
# Try this:  
vignette("dplyr")
```

Help

Also try `vignette()` for a more detailed introduction to many packages

```
# Try this:  
vignette("dplyr")
```

Vignettes are a very easy way to learn how and when to use a package

What are objects?

We won't go into OOP details but here are some objects that we'll be working with regularly:

- vectors
- matrices
- data frames
- lists
- functions
- etc.

What are objects?

We won't go into OOP details but here are some objects that we'll be working with regularly:

- vectors
- matrices
- data frames
- lists
- functions
- etc.

A lot of these are probably familiar if you have coding experience

What are objects?

We won't go into OOP details but here are some objects that we'll be working with regularly:

- vectors
- matrices
- data frames
- lists
- functions
- etc.

A lot of these are probably familiar if you have coding experience

But there are always language-specific features/subtleties

Global environment

```
## Create a small data frame called "df".  
df ← data.frame(x = 1:2, y = 3:4)  
df
```

```
##   x y  
## 1 1 3  
## 2 2 4
```

Global environment

```
## Create a small data frame called "df".  
df ← data.frame(x = 1:2, y = 3:4)  
df
```

```
##   x y  
## 1 1 3  
## 2 2 4
```

Now, let's try to run a regression¹ on these "x" and "y" variables:

¹ Yes, this is a dumb regression with perfectly co-linear variables. Just go with it.

Global environment

```
lm(y ~ x) ## The "lm" stands for linear model(s)
```

Global environment

```
lm(y ~ x) ## The "lm" stands for linear model(s)
```

```
## Error in eval(predvars, data, env): object 'y' not found
```

Global environment

```
lm(y ~ x) ## The "lm" stands for linear model(s)
```

```
## Error in eval(predvars, data, env): object 'y' not found
```

Error?

Global environment

The error message is

```
## Error in eval(predvars, data, env): object 'y' not found
```

Global environment

The error message is

```
## Error in eval(predvars, data, env): object 'y' not found
```

R can't find the variables that we've supplied in our **Global Environment**

Global environment

The error message is

```
## Error in eval(predvars, data, env): object 'y' not found
```

R can't find the variables that we've supplied in our **Global Environment**

Can you find x or y in the RStudio panel?

Global environment

We have to tell R `x` and `y` are a part of the object `df`

Global environment

We have to tell R `x` and `y` are a part of the object `df`

How?

Global environment

We have to tell R `x` and `y` are a part of the object `df`

How?

There are various ways to solve this problem. One is to simply specify the datasource:

```
lm(y ~ x, data = df) ## Works when we add "data = df"!
```

```
##  
## Call:  
## lm(formula = y ~ x, data = df)  
##  
## Coefficients:  
## (Intercept)          x  
##               2              1
```

Global environment: why it matters

This matters largely for Stata users

Global environment: why it matters

This matters largely for Stata users

In Stata, the workspace is basically just a single data frame \Rightarrow all variables are in the global environment

Global environment: why it matters

This matters largely for Stata users

In Stata, the workspace is basically just a single data frame ⇒ all variables are in the global environment

Big problem with this is you can't have multiple data frames / datasets in memory

Working with multiple objects

We can create a second data frame in memory!

```
df2 <- data.frame(x = rnorm(10), y = runif(10))  
df
```

```
##   x  y  
## 1 1  3  
## 2 2  4
```

```
df2
```

```
##           x         y  
## 1 -0.613382890 0.2259774  
## 2  0.130896371 0.2824887  
## 3 -0.905843002 0.2304107  
## 4 -1.136636840 0.6197059  
## 5  1.131662783 0.1040578  
## 6 -0.009838234 0.6782971  
## 7 -0.265665242 0.5561162
```

Indexing

How do we index in R?

Indexing

How do we index in R?

We've already seen an example of indexing in the form of R console output:

```
1+2
```

```
## [1] 3
```

The `[1]` above denotes the first (and, in this case, only) element of our output.¹

Indexing

How do we index in R?

We've already seen an example of indexing in the form of R console output:

```
1+2
```

```
## [1] 3
```

The `[1]` above denotes the first (and, in this case, only) element of our output.¹

In this case, a vector of length one equal to the value "3"

Indexing

Try the following in your console to see a more explicit example of indexed output:

```
rnorm(n = 100, mean = 0, sd = 1)
```

```
## [1] 0.30361622 -0.57146408 -0.30867380 -0.40955540 0.04517408 0.10430013
## [7] 0.17380547 -0.95164107 -0.51172648 -0.22425068 1.33458580 1.15499579
## [13] 2.05312828 0.17165531 -0.18688468 0.58631875 0.13221947 -0.28582323
## [19] 0.72600767 0.71954104 0.61187714 -0.59878087 -0.97809484 -1.32639439
## [25] 0.11650066 0.15587925 -1.57710485 2.50351974 -0.37321754 2.17030535
## [31] 0.01364178 0.19554214 0.47548680 0.77800681 -1.82338236 0.06250832
## [37] 1.12761918 0.51172542 -1.13028292 0.66446996 0.19853222 0.53383038
## [43] -0.62655535 -2.46440515 -0.59379352 1.31853678 1.34202967 0.01819441
## [49] 0.64271498 -0.71066687 0.87527370 -0.44845575 -1.13081044 0.56774060
## [55] -0.97787131 -1.03423017 -0.98067760 -0.14674311 0.53411563 0.10658554
## [61] -0.46979377 -0.46079817 -1.01536509 0.82698389 -1.40094348 -0.67181277
## [67] -0.08459309 1.15130433 0.03260988 -0.26004754 -1.31336998 -0.42364185
## [73] 0.85063648 -0.06952117 -0.48511227 -0.77896876 -0.22249120 -0.56392488
## [79] 0.74349573 -0.99275277 -1.78652358 -0.13506962 0.89185660 0.04830713
```

Indexing: []

We can also use [] to index objects that we create in R.

```
a <- 11:20  
a
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
a[4] ## Get the 4th element of object "a"
```

```
## [1] 14
```

```
a[c(4, 6)] ## Get the 4th and 6th elements
```

```
## [1] 14 16
```

Indexing: []

It also works on larger arrays (vectors, matrices, data frames, and lists). For example:

```
starwars[1, 1] ## Show the cell corresponding to the 1st row & 1st column of the data frame.
```

```
## # A tibble: 1 x 1
##   name
##   <chr>
## 1 Luke Skywalker
```

Indexing: []

It also works on larger arrays (vectors, matrices, data frames, and lists). For example:

```
starwars[1, 1] ## Show the cell corresponding to the 1st row & 1st column of the data frame.
```

```
## # A tibble: 1 x 1
##   name
##   <chr>
## 1 Luke Skywalker
```

What does `starwars[1:3, 1]` give you?

Indexing: []

We haven't covered them properly yet, but **lists** are a more complex type of array object in R

Indexing: []

We haven't covered them properly yet, but **lists** are a more complex type of array object in R

They can contain a random assortment of objects that don't share the same characteristics

Indexing: []

We haven't covered them properly yet, but **lists** are a more complex type of array object in R

They can contain a random assortment of objects that don't share the same characteristics

- e.g. a list can contain a scalar, a string, and a data frame, or even another list

Indexing: []

The relevance to indexing is that lists require two square brackets [[]] to index the parent list item and then the standard [] within that parent item:

```
my_list <- list(  
  a = "hello",  
  b = c(1,2,3),  
  c = data.frame(x = 1:5, y = 6:10)  
)  
my_list[[1]] ## Return the 1st list object
```

```
## [1] "hello"
```

```
my_list[[2]][3] ## Return the 3rd element of the 2nd list object
```

```
## [1] 3
```

Indexing: \$

Lists provide a nice segue to our other indexing operator: \$

- Let's continue with the my_list example from the previous slide.

```
my_list
```

```
## $a
## [1] "hello"
##
## $b
## [1] 1 2 3
##
## $c
##   x   y
## 1 1   6
## 2 2   7
## 3 3   8
## 4 4   9
## 5 5  10
```

Indexing: \$

Lists provide a nice segue to our other indexing operator: \$.

- Let's continue with the my_list example

```
my_list
```

```
## $a
## [1] "hello"
##
## $b
## [1] 1 2 3
##
## $c
##   x  y
## 1 1  6
## 2 2  7
## 3 3  8
## 4 4  9
## 5 5 10
```

Indexing: \$

We can call these objects directly by name using the dollar sign, e.g.

```
my_list$a ## Return list object "a"
```

```
## [1] "hello"
```

```
my_list$b[3] ## Return the 3rd element of list object "b"
```

```
## [1] 3
```

```
my_list$c$x ## Return column "x" of list object "c"
```

```
## [1] 1 2 3 4 5
```

Indexing: \$

The `$` form of indexing also works for other object types

In some cases, you can also combine the two index options:

```
starwars$name[1] # first element of the name column of the starwars data frame
```

```
## [1] "Luke Skywalker"
```

Indexing: \$

However, note some key differences between the output from this example and that of our previous `starwars[1, 1]` example:

```
starwars$name[1]
```

```
## [1] "Luke Skywalker"
```

```
starwars[1, 1]
```

```
## # A tibble: 1 x 1
##   name
##   <chr>
## 1 Luke Skywalker
```

Removing objects

Use `rm()` to remove an object or objects from your working environment.

```
a <- "hello"  
b <- "world"  
rm(a, b)
```

You can also use `rm(list = ls())` to remove all objects in your working environment (except packages), but this is **frowned upon**

Removing objects

Use `rm()` to remove an object or objects from your working environment.

```
a <- "hello"  
b <- "world"  
rm(a, b)
```

You can also use `rm(list = ls())` to remove all objects in your working environment (except packages), but this is **frowned upon**

Just start a new R session instead

The tidyverse

What is "tidy" data?

Resources:

- [Vignette](#) (from the `tidyr` package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

What is "tidy" data?

Resources:

- [Vignette](#) (from the `tidyr` package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

Key points:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

What is "tidy" data?

Resources:

- [Vignette](#) (from the `tidyr` package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

Key points:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

Basically, tidy data is more likely to be [long \(i.e. narrow\)](#) than wide

Checklist

Install tidyverse: `install.packages('tidyverse')`

Install nycflights13: `install.packages('nycflights13', repos = 'https://cran.rstudio.com')`

Tidyverse vs. base R

Lots of debate over tidyverse vs base R

Tidyverse vs. base R

Lots of debate over tidyverse vs base R

The answer is **obvious**: We should teach the tidyverse first

- Good documentation and support
- Consistent philosophy and syntax
- Nice front-end for big data tools
- For data cleaning, plotting, the tidyverse is elite

Tidyverse vs. base R

Base R is still great

- Base R is extremely flexible and powerful
- The tidyverse can't do everything
- Using base R and the tidyverse together is often a good idea

Tidyverse vs. base R

One point of convenience is that there is often a direct correspondence between a tidyverse command and its base R equivalent:

tidyverse	base
?readr::read_csv	?utils::read.csv
?dplyr::if_else	?base::ifelse
?tibble::tibble	?base :: data.frame

Tidyverse functions typically have extra features on top of base R

Tidyverse vs. base R

One point of convenience is that there is often a direct correspondence between a tidyverse command and its base R equivalent:

tidyverse	base
?readr::read_csv	?utils::read.csv
?dplyr::if_else	?base::ifelse
?tibble::tibble	?base :: data.frame

Tidyverse functions typically have extra features on top of base R

There are always many ways to achieve a single goal in R

Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

We have actually loaded a number of packages: **ggplot2**, **tibble**, **dplyr**, etc

Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

We have actually loaded a number of packages: **ggplot2**, **tibble**, **dplyr**, etc

We can also see information about the package versions and some
namespace conflicts

Tidyverse packages

The tidyverse actually comes with a lot more packages than those that are just loaded automatically

```
tidyverse_packages()
```

```
## [1] "broom"        "cli"          "crayon"       "dbplyr"       "dplyr"  
## [6] "forcats"      "ggplot2"       "haven"        "hms"         "httr"  
## [11] "jsonlite"     "lubridate"    "magrittr"     "modelr"      "pillar"  
## [16] "purrr"        "readr"        "readxl"       "reprex"      "rlang"  
## [21] "rstudioapi"   "rvest"        "stringr"     "tibble"      "tidyverse"  
## [26] "xml2"         "tidyverse"
```

e.g. the **lubridate** package is for working with dates and the **rvest** package is for webscraping

Tidyverse packages

The tidyverse actually comes with a lot more packages than those that are just loaded automatically

```
tidyverse_packages()
```

```
## [1] "broom"        "cli"          "crayon"       "dbplyr"       "dplyr"  
## [6] "forcats"      "ggplot2"       "haven"        "hms"         "httr"  
## [11] "jsonlite"     "lubridate"    "magrittr"     "modelr"      "pillar"  
## [16] "purrr"        "readr"        "readxl"       "reprex"      "rlang"  
## [21] "rstudioapi"   "rvest"        "stringr"     "tibble"      "tidyverse"  
## [26] "xml2"         "tidyverse"
```

e.g. the **lubridate** package is for working with dates and the **rvest** package is for webscraping

These packages have to be loaded separately

Tidyverse packages

We're going to focus on two workhorse packages:

1. **dplyr**
2. **tidyr**

These are the packages for cleaning and wrangling data

Tidyverse packages

We're going to focus on two workhorse packages:

1. **dplyr**
2. **tidyr**

These are the packages for cleaning and wrangling data

They are thus the ones that you will likely make the most use of

Tidyverse packages

We're going to focus on two workhorse packages:

1. **dplyr**
2. **tidyr**

These are the packages for cleaning and wrangling data

They are thus the ones that you will likely make the most use of

Data cleaning and wrangling occupies an inordinate amount of time, no matter where you are in your research career

Pipes: %>%

The pipe operator `%>%` lets us perform a sequence of operations in a very nice and tidy way

Pipes: %>%

The pipe operator `%>%` lets us perform a sequence of operations in a very nice and tidy way

Suppose we wanted to figure out the average highway miles per gallon of Audi's in the `mpg` dataset:

```
mpg
```

```
## # A tibble: 234 x 11
##   manufacturer model     displ  year   cyl trans  drv   cty   hwy fl class
##   <chr>        <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi         a4      1.8  1999     4 auto(l... f       18    29 p     comp...
## 2 audi         a4      1.8  1999     4 manual... f      21    29 p     comp...
## 3 audi         a4      2.0  2008     4 manual... f      20    31 p     comp...
## 4 audi         a4      2.0  2008     4 auto(a... f      21    30 p     comp...
## 5 audi         a4      2.8  1999     6 auto(l... f      16    26 p     comp...
## 6 audi         a4      2.8  1999     6 manual... f      18    26 p     comp...
```

Pipes: %>%

There's two ways you might do this without taking advantage of pipes:

Pipes: %>%

There's two ways you might do this without taking advantage of pipes:

The first is to do it step-by-step, line-by-line which requires a lot of variable assignment

```
audis_mpg <- filter(mpg, manufacturer == "audi")
audis_mpg_grouped <- group_by(filter(mpg, manufacturer == "audi"), model)
summarise(audis_mpg_grouped, hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>       <dbl>
## 1 a4          28.3
## 2 a4 quattro  25.8
## 3 a6 quattro  24
```

Pipes: %>%

Next you could do it all in one line which is hard to read

```
summarise(group_by(filter(mpg, manufacturer="audi")), model), hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>       <dbl>
## 1 a4          28.3
## 2 a4 quattro  25.8
## 3 a6 quattro  24
```

Pipes: %>%

Or, you could use **pipes** %>%:

```
mpg %>% filter(manufacturer="audi") %>% group_by(model) %>% summarise(hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>       <dbl>
## 1 a4          28.3
## 2 a4 quattro  25.8
## 3 a6 quattro  24
```

Pipes: %>%

Or, you could use **pipes** %>%:

```
mpg %>% filter(manufacturer="audi") %>% group_by(model) %>% summarise(hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>       <dbl>
## 1 a4          28.3
## 2 a4 quattro  25.8
## 3 a6 quattro  24
```

It performs the operations from left to right, exactly like you'd think of them:
take this object (mpg), do this (filter), then do this (group by car model), then
do this (take the mean of highway miles)

Use vertical space

Pipes are even more readable if we write it over several lines:

```
mpg %>%
  filter(manufacturer=="audi") %>%
  group_by(model) %>%
  summarise(hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>        <dbl>
## 1 a4            28.3
## 2 a4 quattro    25.8
## 3 a6 quattro    24
```

Using vertical space costs nothing and makes for much more readable code

dplyr

Aside: dplyr 1.0.0 release

Please make sure that you are running at least **dplyr** 1.0.0 before continuing.

```
packageVersion('dplyr')
```

```
## [1] '1.0.5'
```

```
# install.packages('dplyr') ## install updated version if < 1.0.0
```

The five key dplyr verbs

1. `filter`: Subset/filter rows based on their values
2. `arrange`: Reorder/arrange rows based on their values
3. `select`: Select columns/variables
4. `mutate`: Create new columns/variables
5. `summarise`: Collapse multiple rows into a single summary value, potentially by a grouping variable

The five key dplyr verbs

1. `filter`: Subset/filter rows based on their values
2. `arrange`: Reorder/arrange rows based on their values
3. `select`: Select columns/variables
4. `mutate`: Create new columns/variables
5. `summarise`: Collapse multiple rows into a single summary value, potentially by a grouping variable

Let's practice these commands together using the `starwars` data frame that comes pre-packaged with `dplyr`

Starwars

Here's the `starwars` dataset, it has 87 observations of 14 variables

```
starwars
```

```
## # A tibble: 87 x 14
##   name    height  mass hair_color skin_color eye_color birth_year sex gender
##   <chr>     <int> <dbl> <chr>       <chr>      <chr>        <dbl> <chr> <chr>
## 1 Luke Skywalker 172     77  blond       fair        blue          19  male  masculin...
## 2 C-3PO            167     75 <NA>        gold        yellow        112  none  masculin...
## 3 R2-D2             96     32 <NA>        white, bl... red           33  none  masculin...
## 4 Darth Vader     202    136  none        white        yellow        41.9 male  masculin...
## 5 Leia Organa     150     49  brown       light       brown          19  femal... feminin...
## 6 Owen Lars       178    120  brown, grey light       blue          52  male  masculin...
## 7 Beru Whitesun  165     75  brown       light       blue           47  femal... feminin...
## 8 R5-D4            97     32 <NA>        white, red red           NA  none  masculin...
## 9 Biggs Darko     183     84  black       light       brown          24  male  masculin...
## 10 Obi-Wan Kenobi 182     77  auburn, wh... fair        blue-gray       57  male  masculin...
## # ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,
## #   films <list>, vehicles <list>, starships <list>
```

1) dplyr::filter

Here we are subsetting the observations of humans that are at least 190cm

```
starwars %>%  
  filter(  
    species = "Human",  
    height >= 190  
)
```

```
## # A tibble: 4 x 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex   gender  
##   <chr>     <int> <dbl> <chr>       <chr>       <chr>        <dbl> <chr> <chr>  
## 1 Darth Va...     202    136 none       white       yellow        41.9 male   masculin...  
## 2 Qui-Gon ...     193     89 brown      fair        blue         92   male   masculin...  
## 3 Dooku        193     80 white      fair        brown        102  male   masculin...  
## 4 Bail Pre...     191     NA black      tan         brown        67   male   masculin...  
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

1) dplyr::filter

You can filter using regular expressions with grep-type commands or the `stringr` package

```
starwars %>%  
  filter(stringr::str_detect(name, "Skywalker"))
```



```
## # A tibble: 3 x 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex gender  
##   <chr>     <int> <dbl> <chr>       <chr>       <chr>       <dbl> <chr> <chr>  
## 1 Luke Sk...     172     77 blond      fair        blue         19 male   masculin...  
## 2 Anakin ...    188     84 blond      fair        blue        41.9 male   masculin...  
## 3 Shmi Sk...    163     NA black      fair        brown        72 female feminin...  
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

This subsets the observations for individuals whose names contain "Skywalker"

1) dplyr::filter

A very common `filter` use case is identifying/removing missing data cases:

```
starwars %>%  
  filter(is.na(height))  
  
## # A tibble: 6 x 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex    gender  
##   <chr>     <int> <dbl> <chr>       <chr>       <dbl> <chr> <chr>  
## 1 Arvel C...     NA     NA brown       fair        brown         NA male  masculin...  
## 2 Finn          NA     NA black       dark        dark          NA male  masculin...  
## 3 Rey           NA     NA brown       light       hazel         NA female feminin...  
## 4 Poe Dam...     NA     NA brown       light       brown         NA male  masculin...  
## 5 BB8            NA     NA none        none       black         NA none  masculin...  
## 6 Captain...     NA     NA unknown     unknown     unknown        NA <NA>  <NA>  
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

1) dplyr::filter

To remove missing observations, use negation:

```
starwars %>%  
  filter(!is.na(height))
```

```
## # A tibble: 81 x 14  
##   name    height  mass hair_color  skin_color eye_color birth_year sex gender  
##   <chr>     <int> <dbl> <chr>       <chr>      <chr>          <dbl> <chr> <chr>  
## 1 Luke Skywalker 172     77 blond       fair        blue            19 male  masculin...  
## 2 C-3PO           167     75 <NA>        gold        yellow         112 none  masculin...  
## 3 R2-D2           96      32 <NA>       white, bl... red             33 none  masculin...  
## 4 Darth Vader 202     136 none        white        yellow         41.9 male  masculin...  
## 5 Leia Organa 150      49 brown       light        brown            19 female feminin...  
## 6 Owen Lars 178     120 brown, grey light        blue            52 male  masculin...  
## 7 Beru Whitesun 165     75 brown       light        blue            47 female feminin...  
## 8 R5-D4           97      32 <NA>       white, red red             NA none  masculin...  
## 9 Biggs Darko 183      84 black       light        brown            24 male  masculin...  
## 10 Obi-Wan Kenobi 182     77 auburn, wh... fair        blue-gray         57 male  masculin...  
## # ... with 71 more rows, and 5 more variables: homeworld <chr>, species <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

2) dplyr::arrange

`arrange` sorts the data frame based on the variables you supply:

```
starwars %>%  
  arrange(birth_year)  
  
## # A tibble: 87 x 14  
##   name    height  mass hair_color skin_color eye_color birth_year sex gender  
##   <chr>     <int> <dbl> <chr>      <chr>      <chr>        <dbl> <chr> <chr>  
## 1 Wicket ...     88    20  brown      brown      brown            8 male  masculin...  
## 2 IG-88          200   140 none       metal      red              15 none  masculin...  
## 3 Luke Sk...     172    77 blond      fair       blue             19 male  masculin...  
## 4 Leia Or...     150    49 brown      light      brown            19 female feminin...  
## 5 Wedge A...     170    77 brown      fair       hazel            21 male  masculin...  
## 6 Plo Koon       188    80 none       orange     black            22 male  masculin...  
## 7 Biggs D...     183    84 black      light      brown            24 male  masculin...  
## 8 Han Solo       180    80 brown      fair       brown            29 male  masculin...  
## 9 Lando C...     177    79 black      dark       brown            31 male  masculin...  
## 10 Boba Fe...    183   78.2 black     fair       brown           31.5 male  masculin...  
## # ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

2) dplyr::arrange

We can also arrange items in descending order using `arrange(desc())`

```
starwars %>%  
  arrange(desc(birth_year))
```

```
## # A tibble: 87 x 14  
##   name    height  mass hair_color skin_color eye_color birth_year sex gender  
##   <chr>     <int> <dbl> <chr>      <chr>      <chr>        <dbl> <chr> <chr>  
## 1 Yoda       66     17  white      green       brown         896 male  masculin...  
## 2 Jabba     ... 175    1358 <NA>      green-tan,... orange        600 herm... masculin...  
## 3 Chewba...  228    112  brown      unknown      blue          200 male  masculin...  
## 4 C-3PO      167     75 <NA>      gold        yellow        112 none  masculin...  
## 5 Dooku      193     80  white      fair        brown         102 male  masculin...  
## 6 Qui-Go...  193     89  brown      fair        blue          92 male  masculin...  
## 7 Ki-Adi...  198     82  white      pale        yellow        92 male  masculin...  
## 8 Finis ...  170     NA  blond      fair        blue          91 male  masculin...  
## 9 Palpat...  170     75  grey       pale        yellow        82 male  masculin...  
## 10 Cliegg... 183     NA  brown      fair        blue          82 male  masculin...  
## # ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

3) dplyr::select

Use commas to select multiple columns out of a data frame, deselect a column with "-", select across multiple columns with "first:last":

```
starwars %>%  
  select(name:skin_color, species, -height)
```

```
## # A tibble: 87 x 5  
##   name           mass hair_color   skin_color species  
##   <chr>        <dbl> <chr>       <chr>     <chr>  
## 1 Luke Skywalker    77  blond      fair       Human  
## 2 C-3PO              75 <NA>       gold      Droid  
## 3 R2-D2              32 <NA>       white, blue Droid  
## 4 Darth Vader       136 none       white      Human  
## 5 Leia Organa        49 brown      light      Human  
## 6 Owen Lars          120 brown, grey light      Human  
## 7 Beru Whitesun lars 75 brown      light      Human  
## 8 R5-D4              32 <NA>       white, red Droid  
## 9 Biggs Darklighter  84 black      light      Human  
## 10 Obi-Wan Kenobi    77 auburn, white fair      Human
```

3) dplyr::select

You can also rename your selected variables in place

```
starwars %>%  
  select(alias = name, crib = homeworld)
```

```
## # A tibble: 87 x 2  
##   alias      crib  
##   <chr>      <chr>  
## 1 Luke Skywalker Tatooine  
## 2 C-3PO        Tatooine  
## 3 R2-D2        Naboo  
## 4 Darth Vader Tatooine  
## 5 Leia Organa Alderaan  
## 6 Owen Lars   Tatooine  
## 7 Beru Whitesun lars Tatooine  
## 8 R5-D4        Tatooine  
## 9 Biggs Darklighter Tatooine  
## 10 Obi-Wan Kenobi Stewjon  
## # ... with 77 more rows
```

3) dplyr::select

If you just want to rename columns without subsetting them, you can use

rename:

```
starwars %>%  
  rename(alias = name, crib = homeworld)  
  
## # A tibble: 87 x 14  
##   alias    height  mass hair_color  skin_color eye_color birth_year sex gender  
##   <chr>     <int> <dbl> <chr>       <chr>      <chr>        <dbl> <chr> <chr>  
## 1 Luke S...     172     77 blond      fair       blue          19 male  masculin...  
## 2 C-3PO         167     75 <NA>       gold       yellow        112 none  masculin...  
## 3 R2-D2          96     32 <NA>       white, bl... red          33 none  masculin...  
## 4 Darth ...      202    136 none       white       yellow        41.9 male  masculin...  
## 5 Leia O...      150     49 brown      light      brown          19 femal... feminin...  
## 6 Owen L...      178    120 brown, grey light      blue          52 male  masculin...  
## 7 Beru W...      165     75 brown      light      blue          47 femal... feminin...  
## 8 R5-D4          97     32 <NA>       white, red red           NA none  masculin...  
## 9 Biggs ...      183     84 black      light      brown          24 male  masculin...  
## 10 Obi-Wa       182     77 auburn wh... fair      blue-gray        57 male  masculin...
```

3) dplyr::select cont.

The `select(contains(PATTERN))` option provides a nice shortcut in relevant cases.

```
starwars %>%  
  select(name, contains("color"))  
  
## # A tibble: 87 x 4  
##   name           hair_color   skin_color eye_color  
##   <chr>          <chr>       <chr>      <chr>  
## 1 Luke Skywalker    blond      fair       blue  
## 2 C-3PO              <NA>       gold       yellow  
## 3 R2-D2              <NA>       white, blue red  
## 4 Darth Vader        none       white      yellow  
## 5 Leia Organa         brown      light      brown  
## 6 Owen Lars           brown, grey light      blue  
## 7 Beru Whitesun lars brown      light      blue  
## 8 R5-D4              <NA>       white, red red  
## 9 Biggs Darklighter  black      light      brown  
## 10 Obi-Wan Kenobi     auburn    white, fair blue-gray
```

3) dplyr::select

The `select(... , everything())` option is another useful shortcut if you only want to bring some variable(s) to the "front" of a data frame

```
starwars %>%
  select(species, homeworld, everything()) %>%
  head(5)

## # A tibble: 5 x 14
##   species homeworld name           height  mass hair_color skin_color eye_color
##   <chr>     <chr>    <chr>        <int>  <dbl>   <chr>       <chr>      <chr>
## 1 Human     Tatooine  Luke Skywalker    172     77  blond       fair       blue
## 2 Droid      Tatooine  C-3PO          167     75 <NA>        gold       yellow
## 3 Droid      Naboo     R2-D2           96      32 <NA>      white, blue red
## 4 Human     Tatooine  Darth Vader     202    136  none        white      yellow
## 5 Human     Alderaan  Leia Organa     150     49  brown       light      brown
## # ... with 6 more variables: birth_year <dbl>, sex <chr>, gender <chr>,
## #   films <list>, vehicles <list>, starships <list>
```

3) dplyr::select

You can also use `relocate` to do the same thing

```
starwars %>%  
  relocate(species, homeworld) %>%  
  head(5)
```



```
## # A tibble: 5 x 14  
##   species homeworld name           height  mass hair_color skin_color eye_color  
##   <chr>     <chr>    <chr>        <int>  <dbl> <chr>       <chr>      <chr>  
## 1 Human     Tatooine  Luke Skywalker     172     77  blond       fair       blue  
## 2 Droid      Tatooine  C-3PO          167     75 <NA>        gold       yellow  
## 3 Droid      Naboo     R2-D2          96      32 <NA>        white, blue red  
## 4 Human     Tatooine  Darth Vader      202    136  none        white       yellow  
## 5 Human     Alderaan  Leia Organa      150     49  brown       light      brown  
## # ... with 6 more variables: birth_year <dbl>, sex <chr>, gender <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

4) dplyr::mutate

You can create new columns from scratch as transformations of existing columns:

```
starwars %>%  
  select(name, birth_year) %>%  
  mutate(dog_years = birth_year * 7) %>%  
  mutate(comment = paste0(name, " is ", dog_years, " in dog years.))
```

```
## # A tibble: 87 x 4  
##   name           birth_year  dog_years comment  
##   <chr>          <dbl>      <dbl> <chr>  
## 1 Luke Skywalker     19        133  Luke Skywalker is 133 in dog years.  
## 2 C-3PO              112       784  C-3PO is 784 in dog years.  
## 3 R2-D2              33        231  R2-D2 is 231 in dog years.  
## 4 Darth Vader        41.9      293. Darth Vader is 293.3 in dog years.  
## 5 Leia Organa         19        133  Leia Organa is 133 in dog years.  
## 6 Owen Lars            52       364  Owen Lars is 364 in dog years.  
## 7 Beru Whitesun lars    47       329  Beru Whitesun lars is 329 in dog yea...  
## 8 R5-D4              NA        NA   R5-D4 is NA in dog years
```

4) dplyr::mutate

Note: `mutate` creates variables in order, so you can chain multiple mutates in a single call

```
starwars %>%  
  select(name, birth_year) %>%  
  mutate(  
    dog_years = birth_year * 7, ## Separate with a comma  
    comment = paste0(name, " is ", dog_years, " in dog years.")  
  )
```

```
## # A tibble: 87 x 4  
##   name      birth_year  dog_years comment  
##   <chr>        <dbl>     <dbl> <chr>  
## 1 Luke Skywalker      19       133  Luke Skywalker is 133 in dog years.  
## 2 C-3PO                 112      784  C-3PO is 784 in dog years.  
## 3 R2-D2                  33      231  R2-D2 is 231 in dog years.  
## 4 Darth Vader             41.9    293. Darth Vader is 293.3 in dog years.  
## 5 Leia Organa              19       133  Leia Organa is 133 in dog years.  
## 6 Owen Lars                52      364  Owen Lars is 364 in dog years.
```

4) dplyr::mutate

Boolean, logical and conditional operators all work well with `mutate` too:

```
starwars %>%
  select(name, height) %>%
  filter(name %in% c("Luke Skywalker", "Anakin Skywalker")) %>%
  mutate(tall1 = height > 180) %>% # TRUE or FALSE
  mutate(tall2 = ifelse(height > 180, "Tall", "Short")) ## Same effect, but can choose labels
```

```
## # A tibble: 2 x 4
##   name           height tall1 tall2
##   <chr>          <int> <lgl> <chr>
## 1 Luke Skywalker     172 FALSE Short
## 2 Anakin Skywalker    188 TRUE  Tall
```

4) dplyr::mutate

Lastly, combining `mutate` with `across` allows you to easily work on a subset of variables:

```
starwars %>%
  select(name:eye_color) %>%
  mutate(across(where(is.character), toupper)) %>% # Take all character variables, uppercase them
  head(5)
```

```
## # A tibble: 5 x 6
##   name           height  mass hair_color skin_color eye_color
##   <chr>        <int> <dbl> <chr>      <chr>      <chr>
## 1 LUKE SKYWALKER     172    77 BLOND      FAIR       BLUE
## 2 C-3PO              167    75 <NA>       GOLD       YELLOW
## 3 R2-D2               96    32 <NA>      WHITE, BLUE  RED
## 4 DARTH VADER        202   136 NONE       WHITE      YELLOW
## 5 LEIA ORGANA         150    49 BROWN     LIGHT      BROWN
```

5) dplyr::summarise

Summarising useful in combination with the `group_by` command

```
starwars %>%  
  group_by(species, gender) %>% # for each species-gender combo  
  summarise(mean_height = mean(height, na.rm = TRUE)) # calculate the mean height
```

```
## # A tibble: 42 x 3  
## # Groups:   species [38]  
##   species   gender   mean_height  
##   <chr>     <chr>       <dbl>  
## 1 Aleena    masculine      79  
## 2 Besalisk  masculine     198  
## 3 Cerean    masculine     198  
## 4 Chagrian  masculine     196  
## 5 Clawdite  feminine      168  
## 6 Droid     feminine       96  
## 7 Droid     masculine     140  
## 8 Dug       masculine     112  
## 9 Ewok      masculine      88  
## 10 Geonosian masculine     182
```

5) dplyr::summarise

Note that including "na.rm = TRUE" is usually a good idea with summarise functions, it keeps NAs from propagating to the end result

```
## Probably not what we want
starwars %>%
  summarise(mean_height = mean(height))
```

```
## # A tibble: 1 x 1
##   mean_height
##       <dbl>
## 1        NA
```

5) dplyr::summarise

We can also use `across` within `summarise`:

```
starwars %>%  
  group_by(species) %>% # for each species  
  summarise(across(where(is.numeric), mean, na.rm = T)) %>% # take the mean of all numeric variables  
  head(5)
```

```
## # A tibble: 5 x 4  
##   species    height    mass birth_year  
##   <chr>      <dbl>   <dbl>     <dbl>  
## 1 Aleena       79     15       NaN  
## 2 Besalisk     198    102       NaN  
## 3 Cerean       198     82       92  
## 4 Chagrian     196     NaN       NaN  
## 5 Clawdite     168     55       NaN
```

Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping

- Particularly useful with the `summarise` and `mutate` commands

Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping

- Particularly useful with the `summarise` and `mutate` commands

`slice`: Subset rows by position rather than filtering by values

- E.g. `starwars %>% slice(c(1, 5))`

Other dplyr goodies

`pull`: Extract a column from as a data frame as a vector or scalar

- E.g. `starwars %>% filter(gender=="female") %>% pull(height)`

Other dplyr goodies

`pull`: Extract a column from as a data frame as a vector or scalar

- E.g. `starwars %>% filter(gender=="female") %>% pull(height)`

`count` and `distinct`: Number and isolate unique observations

- E.g. `starwars %>% count(species)`, or `starwars %>% distinct(species)`
- You could also use a combination of `mutate`, `group_by`, and `n()`, e.g.
`starwars %>% group_by(species) %>% mutate(num = n())`.

Other dplyr goodies

There are also a whole class of **window functions** for getting leads and lags, percentiles, cumulative sums, etc.

- See `vignette("window-functions")`.

dplyr::xxxx_join

The last set of commands we need are the `join` commands

dplyr::xxxx_join

The last set of commands we need are the `join` commands

These are the same as `merge` in stata but with a bit more functionality

dplyr::xxxx_join

We merge data with **join operations**:

- `inner_join(df1, df2)`
- `left_join(df1, df2)`
- `right_join(df1, df2)`
- `full_join(df1, df2)`
- `semi_join(df1, df2)`
- `anti_join(df1, df2)`

(You can visualize the operations [here](#))

dplyr::xxxx_join

Lets use the data that comes with the the `nycflights13` package.

```
library(nycflights13)
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>     <int>           <int>     <dbl>     <int>           <int>
## 1 2013     1     1      517             515       2     830           819
## 2 2013     1     1      533             529       4     850           830
## 3 2013     1     1      542             540       2     923           850
## 4 2013     1     1      544             545      -1    1004          1022
## 5 2013     1     1      554             600      -6     812           837
## 6 2013     1     1      554             558      -4     740           728
## 7 2013     1     1      555             600      -5     913           854
## 8 2013     1     1      557             600      -3     709           723
## 9 2013     1     1      557             600      -3     838           846
## 10 2013    1     1      558             600      -2     753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier_labs <fct>, flight_cints <dbl>, tailnum_labs <fct>, origin_labs <fct>, dest_labs <fct>
```

dplyr::xxxx_join

planes

```
## # A tibble: 3,322 x 9
##   tailnum  year type      manufacturer    model engines seats speed engine
##   <chr>    <int> <chr>      <chr>        <chr>    <int> <int> <int> <chr>
## 1 N10156  2004 Fixed wing m... EMBRAER       EMB-1...     2     55   NA Turbo...
## 2 N102UW   1998 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo...
## 3 N103US   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo...
## 4 N104UW   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo...
## 5 N10575   2002 Fixed wing m... EMBRAER       EMB-1...     2     55   NA Turbo...
## 6 N105UW   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo...
## 7 N107US   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo...
## 8 N108UW   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo...
## 9 N109UW   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo...
## 10 N110UW  1999 Fixed wing m... AIRBUS INDUST... A320-...    2    182   NA Turbo...
## # ... with 3,312 more rows
```

Joining operations

Let's perform a left join on the flights and planes datasets

- Note: I'm going to subset columns after the join, but only to keep text on the slide

Joining operations

Let's perform a left join on the flights and planes datasets

- Note: I'm going subset columns after the join, but only to keep text on the slide

```
left_join(flights, planes) %>%
  select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
```

```
## Joining, by = c("year", "tailnum")
```

```
## # A tibble: 336,776 x 10
##   year month   day dep_time arr_time carrier flight tailnum type model
##   <int> <int> <int>    <int>    <int> <chr>    <int> <chr>    <chr> <chr>
## 1  2013     1     1      517      830  UA        1545 N14228  <NA>  <NA>
## 2  2013     1     1      533      850  UA        1714 N24211  <NA>  <NA>
## 3  2013     1     1      542      923  AA        1141 N619AA  <NA>  <NA>
## 4  2013     1     1      544     1004  B6        725  N804JB  <NA>  <NA>
## 5  2013     1     1      554      812  DL        461  N668DN  <NA>  <NA>
```

Joining operations

Note that dplyr made a reasonable guess about which columns to join on (i.e. columns that share the same name), and told us what it chose

```
## Joining, by = c("year", "tailnum")
```

There's an obvious problem here: the variable `year` does not have a consistent meaning across our joining datasets

Joining operations

Note that dplyr made a reasonable guess about which columns to join on (i.e. columns that share the same name), and told us what it chose

```
## Joining, by = c("year", "tailnum")
```

There's an obvious problem here: the variable `year` does not have a consistent meaning across our joining datasets

In one it refers to the *year of flight*, in the other it refers to *year of construction*

Luckily, there's an easy way to avoid this problem: try `?dplyr :: join`

Joining operations

You just need to be more explicit in your join call by using the `by =` argument

```
left_join(  
  flights,  
  planes %>% rename(year_built = year), ## Not necessary w/ below line, but helpful  
  by = "tailnum" ## Be specific about the joining column  
 ) %>%  
 select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, year_built, type, model)  
 head(3) ## Just to save vertical space on the slide
```

```
## # A tibble: 3 x 11  
##   year month   day dep_time arr_time carrier flight tailnum year_built type  
##   <int> <int> <int>    <int>    <int> <chr>    <int> <chr>      <int> <chr>  
## 1  2013     1     1      517      830  UA        1545 N14228      1999 Fixed w...  
## 2  2013     1     1      533      850  UA        1714 N24211      1998 Fixed w...  
## 3  2013     1     1      542      923  AA        1141 N619AA      1990 Fixed w...  
## # ... with 1 more variable: model <chr>
```

Joining operations

Note what happens if we again specify the join column but don't rename the ambiguous `year`:

```
left_join(flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum") %>%
  select(contains("year"), month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
  head(3)

## # A tibble: 3 x 11
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type   model
##   <int>  <int> <int> <int>    <int>    <int> <chr>   <int> <chr>   <chr> <chr>
## 1  2013    1999     1     1      517      830  UA       1545 N14228  Fixe... 737-...
## 2  2013    1998     1     1      533      850  UA       1714 N24211  Fixe... 737-...
## 3  2013    1990     1     1      542      923  AA       1141 N619AA  Fixe... 757-...
```

Joining operations

Note what happens if we again specify the join column but don't rename the ambiguous `year`:

```
left_join(flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum") %>%
  select(contains("year"), month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
  head(3)

## # A tibble: 3 x 11
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type   model
##   <int>  <int> <int> <int>    <int>    <int> <chr>   <int> <chr>   <chr> <chr>
## 1  2013    1999     1     1      517      830  UA        1545 N14228  Fixe... 737-...
## 2  2013    1998     1     1      533      850  UA        1714 N24211  Fixe... 737-...
## 3  2013    1990     1     1      542      923  AA        1141 N619AA  Fixe... 757-...
```

Make sure you know what "year.x" and "year.y" are

tidyr

Key `tidyverse` verbs

1. `pivot_longer`: Pivot wide data into long format (i.e. "melt", "reshape long")
2. `pivot_wider`: Pivot long data into wide format (i.e. "cast", "reshape wide")
3. `separate`: Split one column into multiple columns
4. `unite`: Combine multiple columns into one

Key tidy verbs

1. `pivot_longer`: Pivot wide data into long format (i.e. "melt", "reshape long")
2. `pivot_wider`: Pivot long data into wide format (i.e. "cast", "reshape wide")
3. `separate`: Split one column into multiple columns
4. `unite`: Combine multiple columns into one

Let's practice these verbs together in class

1) `tidyr::pivot_longer`

```
stocks ← data.frame(  
  time = as.Date('2009-01-01') + 0:1,  
  X = rnorm(2, 0, 1),  
  Y = rnorm(2, 0, 2),  
  Z = rnorm(2, 0, 4)  
)  
stocks
```

```
##           time         X         Y         Z  
## 1 2009-01-01  0.6163028 1.779204 -1.461802  
## 2 2009-01-02 -0.1024888 2.548539 -2.816719
```

We have 4 variables, the date and the stocks

How do we get this in tidy form?

1) tidyverse::pivot_longer

```
stocks %>% pivot_longer(-time, names_to = "stock", values_to = "price")
```

We need to pivot the stock name variables x, Y, z longer

1. Choose non-time variables: -time
2. Decide what variable holds the names: names_to = "stock"
3. Decide what variable holds the values: values_to = "price"

1) tidyverse::pivot_longer

```
stocks %>% pivot_longer(-time, names_to = "stock", values_to = "price")
```

```
## # A tibble: 6 x 3
##   time     stock  price
##   <date>    <chr>  <dbl>
## 1 2009-01-01 X      0.616
## 2 2009-01-01 Y      1.78
## 3 2009-01-01 Z     -1.46
## 4 2009-01-02 X     -0.102
## 5 2009-01-02 Y      2.55
## 6 2009-01-02 Z     -2.82
```

1) `tidy::pivot_longer`

Let's quickly save the "tidy" (i.e. long) stocks data frame for use on the next slide

```
tidy_stocks ← stocks %>%  
  pivot_longer(-time, names_to = "stock", values_to = "price")
```

2) tidy::pivot_wider

```
tidy_stocks %>% pivot_wider(names_from = stock, values_from = price)
```

```
## # A tibble: 2 x 4
##   time           X     Y     Z
##   <date>     <dbl> <dbl> <dbl>
## 1 2009-01-01  0.616  1.78 -1.46
## 2 2009-01-02 -0.102  2.55 -2.82
```

```
tidy_stocks %>% pivot_wider(names_from = time, values_from = price)
```

```
## # A tibble: 3 x 3
##   stock `2009-01-01` `2009-01-02`
##   <chr>     <dbl>     <dbl>
## 1 X         0.616    -0.102
## 2 Y         1.78      2.55
## 3 Z        -1.46     -2.82
```

2) tidyverse::pivot_wider

```
tidy_stocks %>% pivot_wider(names_from = stock, values_from = price)
```

```
## # A tibble: 2 x 4
##   time           X     Y     Z
##   <date>     <dbl> <dbl> <dbl>
## 1 2009-01-01  0.616  1.78 -1.46
## 2 2009-01-02 -0.102  2.55 -2.82
```

```
tidy_stocks %>% pivot_wider(names_from = time, values_from = price)
```

```
## # A tibble: 3 x 3
##   stock `2009-01-01` `2009-01-02`
##   <chr>     <dbl>     <dbl>
## 1 X         0.616    -0.102
## 2 Y         1.78      2.55
## 3 Z        -1.46     -2.82
```

Note that the second example has effectively transposed the data

3) tidyverse::separate

```
economists <- data.frame(name = c("Adam.Smith", "Paul.Samuelson", "Milton.Friedman"))
economists
```

```
##           name
## 1      Adam.Smith
## 2  Paul.Samuelson
## 3 Milton.Friedman
```

```
economists %>% separate(name, c("first_name", "last_name"))
```

```
##   first_name last_name
## 1      Adam     Smith
## 2      Paul Samuelson
## 3    Milton   Friedman
```

3) tidyverse::separate

```
conomists <- data.frame(name = c("Adam.Smith", "Paul.Samuelson", "Milton.Friedman"))
conomists
```

```
##           name
## 1      Adam.Smith
## 2  Paul.Samuelson
## 3 Milton.Friedman
```

```
conomists %>% separate(name, c("first_name", "last_name"))
```

```
##   first_name last_name
## 1      Adam     Smith
## 2      Paul Samuelson
## 3    Milton  Friedman
```

This command is pretty smart. But to avoid ambiguity, you can also specify the separation character with `separate(... , sep=".")`

3) `tidyverse::separate`

A related function is `separate_rows`, for splitting up cells that contain multiple fields or observations (a frustratingly common occurrence with survey data)

```
jobs ← data.frame(  
  name = c("Jack", "Jill"),  
  occupation = c("Homemaker", "Philosopher, Philanthropist, Troublemaker")  
)  
jobs
```

```
##      name          occupation  
## 1 Jack           Homemaker  
## 2 Jill Philosopher, Philanthropist, Troublemaker
```

3) `tidyr::separate`

A related function is `separate_rows`, for splitting up cells that contain multiple fields or observations (a frustratingly common occurrence with survey data)

```
## Now split out Jill's various occupations into different rows
jobs %>% separate_rows(occupation)
```

```
## # A tibble: 4 x 2
##   name  occupation
##   <chr> <chr>
## 1 Jack   Homemaker
## 2 Jill   Philosopher
## 3 Jill   Philanthropist
## 4 Jill   Troublemaker
```

4) tidyverse

```
gdp ← data.frame(  
  yr = rep(2016, times = 4),  
  mnth = rep(1, times = 4),  
  dy = 1:4,  
  gdp = rnorm(4, mean = 100, sd = 2)  
)  
gdp
```

```
##      yr mnth dy      gdp  
## 1 2016    1  1 104.73167  
## 2 2016    1  2 102.67609  
## 3 2016    1  3  97.35823  
## 4 2016    1  4  97.84164
```

4) tidyverse

```
## Combine "yr", "mnth", and "dy" into one "date" column  
gdp %>% unite(date, c("yr", "mnth", "dy"), sep = "-")
```

```
##           date      gdp  
## 1 2016-1-1 104.73167  
## 2 2016-1-2 102.67609  
## 3 2016-1-3  97.35823  
## 4 2016-1-4  97.84164
```

4) tidyverse::unite

Note that `unite` will automatically create a character variable:

```
gdp_u <- gdp %>% unite(date, c("yr", "mnth", "dy"), sep = "-") %>% as_tibble()  
gdp_u
```

```
## # A tibble: 4 x 2  
##   date      gdp  
##   <chr>     <dbl>  
## 1 2016-1-1 105.  
## 2 2016-1-2 103.  
## 3 2016-1-3  97.4  
## 4 2016-1-4  97.8
```

4) tidyverse::unite

Note that `unite` will automatically create a character variable:

```
gdp_u <- gdp %>% unite(date, c("yr", "mnth", "dy"), sep = "-") %>% as_tibble()  
gdp_u
```

```
## # A tibble: 4 x 2  
##   date      gdp  
##   <chr>     <dbl>  
## 1 2016-1-1 105.  
## 2 2016-1-2 103.  
## 3 2016-1-3  97.4  
## 4 2016-1-4  97.8
```

If you want to convert it to something else (e.g. date or numeric) then you will need to modify it using `mutate`

4) tidyverse

```
library(lubridate)
gdp_u %>% mutate(date = ymd(date))
```

```
## # A tibble: 4 x 2
##   date      gdp
##   <date>    <dbl>
## 1 2016-01-01 105.
## 2 2016-01-02 103.
## 3 2016-01-03  97.4
## 4 2016-01-04  97.8
```

Other `tidyR` goodies

Use `crossing` to get the full combination of a group of variables

```
crossing(side=c("left", "right"), height=c("top", "bottom"))
```

```
## # A tibble: 4 x 2
##   side  height
##   <chr> <chr>
## 1 left   bottom
## 2 left   top
## 3 right  bottom
## 4 right  top
```

Other tidyverse goodies

Use `crossing` to get the full combination of a group of variables

```
crossing(side=c("left", "right"), height=c("top", "bottom"))
```

```
## # A tibble: 4 x 2
##   side  height
##   <chr> <chr>
## 1 left   bottom
## 2 left   top
## 3 right  bottom
## 4 right  top
```

See `?expand` and `?complete` for more specialized functions that allow you to fill in (implicit) missing data or variable combinations in existing data frames

Randomization and inference

The Rubin causal model

What is a causal effect?

The Rubin causal model

What is a causal effect?

A starting point is that a causal effect is a **comparison between two potential outcomes**

The Rubin causal model

What is a causal effect?

A starting point is that a causal effect is a **comparison between two potential outcomes**

What is the difference in some outcome in the presence vs the absence of a given treatment?

The Rubin causal model

What is a causal effect?

A starting point is that a causal effect is a **comparison between two potential outcomes**

What is the difference in some outcome in the presence vs the absence of a given treatment?

e.g. what is the difference in health between high and low levels of air pollution?

The Rubin causal model

What is a causal effect?

A starting point is that a causal effect is a **comparison between two potential outcomes**

What is the difference in some outcome in the presence vs the absence of a given treatment?

e.g. what is the difference in health between high and low levels of air pollution?

Let's begin formalizing this idea

The Rubin causal model: potential outcomes

Suppose we have a set of **observational units**

- These can be people, states, animals, air quality monitors, etc

Each unit has two **potential outcomes**, but only one is observed

The Rubin causal model: potential outcomes

Suppose we have a set of **observational units**

- These can be people, states, animals, air quality monitors, etc

Each unit has two **potential outcomes**, but only one is observed

The potential outcome is Y_i^1 if unit i received some treatment, and Y_i^0 if the unit did not

The Rubin causal model: potential outcomes

Suppose we have a set of **observational units**

- These can be people, states, animals, air quality monitors, etc

Each unit has two **potential outcomes**, but only one is observed

The potential outcome is Y_i^1 if unit i received some treatment, and Y_i^0 if the unit did not

Y_i^0 corresponds to the control state of the world for i

The Rubin causal model: observable outcomes

Note that these potential outcomes are not the same as observable outcomes

The Rubin causal model: observable outcomes

Note that these potential outcomes are not the same as observable outcomes

Observable outcomes are outcomes that actually show up in the data (i.e. the *factual outcome*)¹

¹The potential outcome that did not happen is called the *counterfactual outcome*.

The Rubin causal model: observable outcomes

Note that these potential outcomes are not the same as observable outcomes

Observable outcomes are outcomes that actually show up in the data (i.e. the *factual* outcome)¹

We can write the observable outcome Y_i as a simple equation:

$$Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$$

where $D_i = 1$ if i received treatment and 0 otherwise

¹The potential outcome that did not happen is called the *counterfactual outcome*.

The Rubin causal model: treatment effects

$$Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$$

This is the basis of the *Rubin causal model*

The Rubin causal model: treatment effects

$$Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$$

This is the basis of the *Rubin causal model*

Rubin defines a treatment/causal effect δ_i as the difference between the two potential outcomes:

$$\delta_i = Y_i^1 - Y_i^0$$

The Rubin causal model: treatment effects

$$Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$$

This is the basis of the *Rubin causal model*

Rubin defines a treatment/causal effect δ_i as the difference between the two potential outcomes:

$$\delta_i = Y_i^1 - Y_i^0$$

This leads to an obvious problem: we only observe one of these two states for each unit i but we need to know both to recover δ_i

The Rubin causal model: average treatment effect

We can derive three parameters of interest from the definition of a treatment effect:

Average treatment effect (ATE):

$$E[\delta_i] = E[Y_i^1 - Y_i^0] = E[Y_i^1] - E[Y_i^0]$$

The Rubin causal model: average treatment effect

We can derive three parameters of interest from the definition of a treatment effect:

Average treatment effect (ATE):

$$E[\delta_i] = E[Y_i^1 - Y_i^0] = E[Y_i^1] - E[Y_i^0]$$

This is the average of the individual treatment effects

The Rubin causal model: average treatment effect

We can derive three parameters of interest from the definition of a treatment effect:

Average treatment effect (ATE):

$$E[\delta_i] = E[Y_i^1 - Y_i^0] = E[Y_i^1] - E[Y_i^0]$$

This is the average of the individual treatment effects

It is also unknowable

Rubin: average treatment on the treated

Average treatment on the treated (ATT):

$$E[\delta_i | D_i = 1] = E[Y_i^1 - Y_i^0 | D_i = 1] = E[Y_i^1 | D_i = 1] - E[Y_i^0 | D_i = 1]$$

Rubin: average treatment on the treated

Average treatment on the treated (ATT):

$$E[\delta_i | D_i = 1] = E[Y_i^1 - Y_i^0 | D_i = 1] = E[Y_i^1 | D_i = 1] - E[Y_i^0 | D_i = 1]$$

This is the average of the individual treatment effects only for the i in the treated group

Rubin: average treatment on the treated

Average treatment on the treated (ATT):

$$E[\delta_i | D_i = 1] = E[Y_i^1 - Y_i^0 | D_i = 1] = E[Y_i^1 | D_i = 1] - E[Y_i^0 | D_i = 1]$$

This is the average of the individual treatment effects only for the i in the treated group

It is also unknowable since we never observe $E[Y_i^0 | D_i = 1]$

Rubin: average treatment on the treated

Average treatment on the treated (ATT):

$$E[\delta_i | D_i = 1] = E[Y_i^1 - Y_i^0 | D_i = 1] = E[Y_i^1 | D_i = 1] - E[Y_i^0 | D_i = 1]$$

This is the average of the individual treatment effects only for the i in the treated group

It is also unknowable since we never observe $E[Y_i^0 | D_i = 1]$

If the treatment effect differs across i then $ATE \neq ATT$

Rubin: average treatment on the untreated

Average treatment on the untreated (ATU):

$$E[\delta_i | D_i = 0] = E[Y_i^1 - Y_i^0 | D_i = 0] = E[Y_i^1 | D_i = 0] - E[Y_i^0 | D_i = 0]$$

Rubin: average treatment on the untreated

Average treatment on the untreated (ATU):

$$E[\delta_i | D_i = 0] = E[Y_i^1 - Y_i^0 | D_i = 0] = E[Y_i^1 | D_i = 0] - E[Y_i^0 | D_i = 0]$$

This is the average of the individual treatment effects only for the i in the untreated group

Rubin: average treatment on the untreated

Average treatment on the untreated (ATU):

$$E[\delta_i | D_i = 0] = E[Y_i^1 - Y_i^0 | D_i = 0] = E[Y_i^1 | D_i = 0] - E[Y_i^0 | D_i = 0]$$

This is the average of the individual treatment effects only for the i in the untreated group

It is also unknowable since we never observe $E[Y_i^1 | D_i = 0]$

Rubin: average treatment on the untreated

Average treatment on the untreated (ATU):

$$E[\delta_i | D_i = 0] = E[Y_i^1 - Y_i^0 | D_i = 0] = E[Y_i^1 | D_i = 0] - E[Y_i^0 | D_i = 0]$$

This is the average of the individual treatment effects only for the i in the untreated group

It is also unknowable since we never observe $E[Y_i^1 | D_i = 0]$

If the treatment effect differs across i then $ATE \neq ATU$

Hands on: understanding treatment effects

We've got our definitions, now lets be a bit more clear as to what we are doing with some empirical examples putting our new R tools to work

Hands on: understanding treatment effects

We've got our definitions, now lets be a bit more clear as to what we are doing with some empirical examples putting our new R tools to work

Suppose treatment D is whether a state has a conservation policy

Hands on: understanding treatment effects

We've got our definitions, now lets be a bit more clear as to what we are doing with some empirical examples putting our new R tools to work

Suppose treatment D is whether a state has a conservation policy

The outcome Y is the number of bird species in that state

Hands on: understanding treatment effects

We've got our definitions, now lets be a bit more clear as to what we are doing with some empirical examples putting our new R tools to work

Suppose treatment D is whether a state has a conservation policy

The outcome Y is the number of bird species in that state

We want to understand the causal effect of conservation policy on the number of species

Hands on: understanding treatment effects

Here's our dataset of both potential outcomes:

```
cons_df # data frame of conservation treatment, outcomes
```

```
## # A tibble: 8 x 4
##   state    Y1     Y0  delta
##   <int> <dbl> <dbl> <dbl>
## 1     1     4     7    -3
## 2     2     7     9    -2
## 3     3     7     0     7
## 4     4    10     1     9
## 5     5     7     7     0
## 6     6     6     0     6
## 7     7     9     3     6
## 8     8    13     4     9
```

Calculate the average treatment effect

Hands on: understanding treatment effects

ATE = $E[\delta_i]$ which we can compute with `dplyr::summarise`:

```
cons_df %>%  
  dplyr::summarise(mean(delta))
```

```
## # A tibble: 1 x 1  
##   `mean(delta)`  
##       <dbl>  
## 1           4
```

The average treatment effect is 4: a conservation policy increases the number of bird species in a state by 4

Hands on: understanding treatment effects

Notice that not all states benefit from conservation policies, and it even backfires in one state

The ATE is just the average over all the different treatment effects

```
cons_df
```

```
## # A tibble: 8 x 4
##   state    Y1     Y0  delta
##   <int> <dbl> <dbl> <dbl>
## 1     1     4     7    -3
## 2     2     7     9    -2
## 3     3     7     0     7
## 4     4    10     1     9
## 5     5     7     7     0
## 6     6     6     0     6
## 7     7     9     3     6
## 8     8    13     4     9
```

Hands on: understanding treatment effects

Suppose we have a perfect policymaker who knows each state's potential outcomes and can perfectly decide whether each state should have a conservation policy

Hands on: understanding treatment effects

Suppose we have a perfect policymaker who knows each state's potential outcomes and can perfectly decide whether each state should have a conservation policy

The policymaker assigns treatment and then observes the actual outcome according to $Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$

Hands on: understanding treatment effects

Suppose we have a perfect policymaker who knows each state's potential outcomes and can perfectly decide whether each state should have a conservation policy

The policymaker assigns treatment and then observes the actual outcome according to $Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$

What does the dataset look like for the observed outcomes?

Hands on: understanding treatment effects

```
observed_df ← cons_df %>%  
  mutate(  
    Y = ifelse(delta > 0, Y1, Y0),  
    D = as.numeric(delta > 0)  
  ) %>%  
  select(state, Y, D)  
observed_df
```

```
## # A tibble: 8 x 3  
##   state     Y     D  
##   <int> <dbl> <dbl>  
## 1     1     7     0  
## 2     2     9     0  
## 3     3     7     1  
## 4     4    10     1  
## 5     5     7     0  
## 6     6     6     1  
## 7     7     9     1  
## 8     8    13     1
```

Hands on: estimating ATEs

Given the **observed** data, what if we tried to *estimate* the ATE by comparing mean outcomes of treated ($D_i = 1$) vs untreated units ($D_i = 0$)

Hands on: estimating ATEs

Given the **observed** data, what if we tried to *estimate* the ATE by comparing mean outcomes of treated ($D_i = 1$) vs untreated units ($D_i = 0$)

This is the simple difference in mean outcomes (SDO):

$$\begin{aligned} SDO &= E[Y^1|D = 1] - E[Y^0|D = 0] \\ &= \frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_U} \sum_{i=1}^{N_U} (y_i | d_i = 0) \end{aligned}$$

where N_T is the number of treated units and N_U is the number of untreated units

Hands on: estimating ATEs

We can compute the SDO using `dplyr::summarise` in conjunction with `dplyr::group_by` on treatment status D :

```
observed_df %>%  
  dplyr::group_by(D) %>%  
  dplyr::summarise(meanY = mean(Y))  
  
## # A tibble: 2 x 2  
##       D   meanY  
##   <dbl> <dbl>  
## 1     0    7.67  
## 2     1     9
```

The SDO is $9 - 7.67 = 1.33 < 4$!

Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Because the SDO is actually composed of three pieces, only one of which is the ATE:¹

$$\text{SDO} = \text{ATE} +$$

See Mixtape pages 89-91 for the derivation.

Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Because the SDO is actually composed of three pieces, only one of which is the ATE:¹

$$SDO = ATE + \text{Selection Bias} +$$

See Mixtape pages 89-91 for the derivation.

Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Because the SDO is actually composed of three pieces, only one of which is the ATE:¹

$$\text{SDO} = \text{ATE} + \text{Selection Bias} + \text{Heterogeneous Treatment Effect Bias}$$

See Mixtape pages 89-91 for the derivation.

Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Because the SDO is actually composed of three pieces, only one of which is the ATE:¹

$$\text{SDO} = \text{ATE} + \text{Selection Bias} + \text{Heterogeneous Treatment Effect Bias}$$

What are these mathematically and intuitively?

See Mixtape pages 89-91 for the derivation.

Hands on: decomposing the SDO

$$\underbrace{E[Y^1|D=1] - E[Y^0|D=0]}_{\text{SDO}} = \underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) + \frac{1}{N_U} \sum_{i=1}^{N_U} (y_i | d_i = 0)}_{\text{SDO}}$$
$$= \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$
$$+ \underbrace{E[Y^0|D=1] - E[Y^0|D=0]}_{\text{Selection bias}}$$
$$+ \underbrace{(1 - \pi)(ATT - ATU)}_{\text{Het. Treat. Eff. Bias}}$$

where π is the share of treated units

Hands on: decomposing the SDO

The LHS we know sums to 1.33 so the RHS must as well

Hands on: decomposing the SDO

The LHS we know sums to 1.33 so the RHS must as well

The first term is the ATE, what we actually want to estimate

Hands on: decomposing the SDO

The LHS we know sums to 1.33 so the RHS must as well

The first term is the ATE, what we actually want to estimate

We **know** ATE = 4 so the last two terms are why SDO < ATE

Hands on: decomposing the SDO

The LHS we know sums to 1.33 so the RHS must as well

The first term is the ATE, what we actually want to estimate

We **know** ATE = 4 so the last two terms are why SDO < ATE

Let's work through these two in more detail

Hands on: Selection bias

The second term is **selection bias**:

$$E[Y^0|D = 1] - E[Y^0|D = 0]$$

It is the inherent differences between the two groups if they did not *actually* get the conservation policy

Hands on: Selection bias

The second term is **selection bias**:

$$E[Y^0|D = 1] - E[Y^0|D = 0]$$

It is the inherent differences between the two groups if they did not *actually* get the conservation policy

how are the two groups different at baseline?

Hands on: Selection bias

The second term is **selection bias**:

$$E[Y^0|D = 1] - E[Y^0|D = 0]$$

It is the inherent differences between the two groups if they did not *actually* get the conservation policy

how are the two groups different at baseline?

The problem is we don't observe $E[Y^0|D = 1]$: what the treated group ($D = 1$) would have looked like without treatment (Y^0)

Hands on: Selection bias

The second term is **selection bias**:

$$E[Y^0|D = 1] - E[Y^0|D = 0]$$

It is the inherent differences between the two groups if they did not *actually* get the conservation policy

how are the two groups different at baseline?

The problem is we don't observe $E[Y^0|D = 1]$: what the treated group ($D = 1$) would have looked like without treatment (Y^0)

Calculate this with the `cons_df` data frame with both potential outcomes

Hands on: Selection bias

```
cons_df %>%
  mutate(
    D = as.numeric(delta > 0)
  ) %>%
  group_by(D) %>%
  summarise(mean(Y0)) # Difference in potential control outcomes across the two groups
```

```
## # A tibble: 2 x 2
##       D `mean(Y0)`
##   <dbl>     <dbl>
## 1     0      7.67
## 2     1      1.6
```

Selection bias is thus $1.6 - 7.67 = -6.07$

Hands on: het. treat. effect bias

The third term is the bias from heterogeneous treatment effects across groups:

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

It is the difference in the effect of the conservation policy across the two groups multiplied by the share that did not get a conservation policy

Hands on: het. treat. effect bias

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

What's the intuition for this?

Hands on: het. treat. effect bias

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

What's the intuition for this?

- How much more of an effect did the policy have on the units that *happened* to get treatment versus the effect the policy would have had on the units that *happened* to be in the control group

Hands on: het. treat. effect bias

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

What's the intuition for this?

- How much more of an effect did the policy have on the units that *happened* to get treatment versus the effect the policy would have had on the units that *happened* to be in the control group
- Scale the size of this difference by fraction of control units

Hands on: het. treat. effect bias

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

What's the intuition for this?

- How much more of an effect did the policy have on the units that *happened* to get treatment versus the effect the policy would have had on the units that *happened* to be in the control group
- Scale the size of this difference by fraction of control units

Calculate this with the `cons_df` data frame with both potential outcomes

Hands on: het. treat. effect bias

```
cons_df %>%
  mutate(
    D = as.numeric(delta > 0)
  ) %>%
  group_by(D) %>%
  summarise(mean(delta), n()) # Difference in potential control and treatment outcomes across the two groups
```

```
## # A tibble: 2 x 3
##       D `mean(delta)` `n()`
##   <dbl>      <dbl>   <int>
## 1     0      -1.67     3
## 2     1       7.4      5
```

Heterogeneous treatment effect bias is thus: $(1 - 5/(5+3)) * (7.4 - (-1.67)) = 3.40$

Hands on: het. treat. effect bias

```
cons_df %>%
  mutate(
    D = as.numeric(delta > 0)
  ) %>%
  group_by(D) %>%
  summarise(mean(delta), n()) # Difference in potential control and treatment outcomes across the two groups
```

```
## # A tibble: 2 x 3
##       D `mean(delta)` `n()`
##   <dbl>      <dbl>   <int>
## 1     0      -1.67     3
## 2     1       7.4      5
```

Heterogeneous treatment effect bias is thus: $(1 - 5/(5+3)) * (7.4 - (-1.67)) = 3.40$

In total we have: **SDO (1.33) = ATE (4) + SB (-6.07) + HTEB (3.40)**

Recap: treatment effect estimates

Taking a simple difference in means of outcomes between treatment and control groups **does** contain what we want

Recap: treatment effect estimates

Taking a simple difference in means of outcomes between treatment and control groups **does** contain what we want

Unfortunately it is **confounded** by two forms of bias:

Recap: treatment effect estimates

Taking a simple difference in means of outcomes between treatment and control groups **does** contain what we want

Unfortunately it is **confounded** by two forms of bias:

1. **Selection bias:** The units in the treatment group different from the control group in the absence of treatment

Recap: treatment effect estimates

Taking a simple difference in means of outcomes between treatment and control groups **does** contain what we want

Unfortunately it is **confounded** by two forms of bias:

1. **Selection bias:** The units in the treatment group different from the control group in the absence of treatment
2. **Heterogeneous treatment effect bias:** The units in the treatment group respond to treatment differently than units in the control group

Bias examples

What are some examples of these forms of bias?

Bias examples

What are some examples of these forms of bias?

Selection bias:

Bias examples

What are some examples of these forms of bias?

Selection bias: If we select certain groups into treatment, for example, if we pass conservation policy in states with little biodiversity

- This may lead to a negative bias on conservation policy impacts, even with a positive ATE

Bias examples

What are some examples of these forms of bias?

Selection bias: If we select certain groups into treatment, for example, if we pass conservation policy in states with little biodiversity

- This may lead to a negative bias on conservation policy impacts, even with a positive ATE

HTEB:

Bias examples

What are some examples of these forms of bias?

Selection bias: If we select certain groups into treatment, for example, if we pass conservation policy in states with little biodiversity

- This may lead to a negative bias on conservation policy impacts, even with a positive ATE

HTEB: If we select units into treatment based on expected response, for example, if we pass policy in states where birds are **very** sensitive to conservation

- This may lead to an overestimate of the size of the treatment effect

Empirical economics

A huge chunk of economics is trying circumvent these forms of bias¹

¹We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

Empirical economics

A huge chunk of economics is trying circumvent these forms of bias¹

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE

¹We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

Empirical economics

A huge chunk of economics is trying circumvent these forms of bias¹

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE
2. Understanding what assumptions are required for these tools to be valid

¹We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

Empirical economics

A huge chunk of economics is trying circumvent these forms of bias¹

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE
2. Understanding what assumptions are required for these tools to be valid
3. Understanding when these ATE estimates are valid

¹We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

Empirical economics

A huge chunk of economics is trying circumvent these forms of bias¹

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE
2. Understanding what assumptions are required for these tools to be valid
3. Understanding when these ATE estimates are valid
4. Seeing how this all works (hands on) in modern environmental economics

¹We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

Empirical economics

A huge chunk of economics is trying circumvent these forms of bias¹

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE
2. Understanding what assumptions are required for these tools to be valid
3. Understanding when these ATE estimates are valid
4. Seeing how this all works (hands on) in modern environmental economics

This should be a less-technical complement to Brian's class

¹We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

Recovering the ATE

There are many different ways to recover the ATE

Recovering the ATE

There are many different ways to recover the ATE

We will cover some subset of:

1. Randomized control trials
2. Regression discontinuity
3. Difference-in-differences
4. Cross-sectional regressions
5. Two way fixed effects

Recovering the ATE

There are many different ways to recover the ATE

We will cover some subset of:

1. Randomized control trials
2. Regression discontinuity
3. Difference-in-differences
4. Cross-sectional regressions
5. Two way fixed effects

All of these approaches have pluses and minuses

Understanding the source of bias

What is the fundamental problem leading to selection bias and HTEB?

Understanding the source of bias

What is the fundamental problem leading to selection bias and HTEB?

Both stem from treatment being correlated with characteristics of the observational units

Understanding the source of bias

What is the fundamental problem leading to selection bias and HTEB?

Both stem from treatment being correlated with characteristics of the observational units

Selection bias is because treatment is correlated with unit characteristics in the absence of treatment

Understanding the source of bias

What is the fundamental problem leading to selection bias and HTEB?

Both stem from treatment being correlated with characteristics of the observational units

Selection bias is because treatment is correlated with unit characteristics in the absence of treatment

HTEB is because treatment is correlated with the size of units' responses to treatment

A potential solution

What's a simple way to deal with this?

A potential solution

What's a simple way to deal with this?

RANDOMIZATION

A potential solution

What's a simple way to deal with this?

RANDOMIZATION

If we randomize treatment across observational units then SDO = ATE

A potential solution

What's a simple way to deal with this?

RANDOMIZATION

If we randomize treatment across observational units then SDO = ATE

Randomization drives selection bias and HTEB to zero

A potential solution

What's a simple way to deal with this?

RANDOMIZATION

If we randomize treatment across observational units then SDO = ATE

Randomization drives selection bias and HTEB to zero

Let's see why

Randomization

Randomization means that treatment will be **independent** of the potential outcomes

Randomization

Randomization means that treatment will be **independent** of the potential outcomes

$$Y^1, Y^0 \perp D$$

Randomization

Randomization means that treatment will be **independent** of the potential outcomes

$$Y^1, Y^0 \perp D$$

In our example we know independence is violated because we assigned the conservation policy to states that had $Y^1 > Y^0$

Randomization

Randomization means that treatment will be **independent** of the potential outcomes

$$Y^1, Y^0 \perp D$$

In our example we know independence is violated because we assigned the conservation policy to states that had $Y^1 > Y^0$

What if we randomized conservation policy?

Randomization

Randomization of treatment / policy means that:

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

Randomization

Randomization of treatment / policy means that:

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

If we randomized treatment across units, then on average, the difference in potential outcomes across groups should be zero

Randomization

Randomization of treatment / policy means that:

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

If we randomized treatment across units, then on average, the difference in potential outcomes across groups should be zero

This means that:

$$E[Y^1|D = 1] = E[Y^1|D = 0] \quad E[Y^0|D = 1] = E[Y^0|D = 0]$$

Randomization

Randomization of treatment / policy means that:

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

If we randomized treatment across units, then on average, the difference in potential outcomes across groups should be zero

This means that:

$$E[Y^1|D = 1] = E[Y^1|D = 0] \quad E[Y^0|D = 1] = E[Y^0|D = 0]$$

The second equation $E[Y^0|D = 1] - E[Y^0|D = 0] = 0$ directly gives us that selection bias is zero with the SDO

Randomization

What does randomization do to HTEB?

Randomization

What does randomization do to HTEB?

$$ATT - ATU$$

$$= (E[Y^1|D = 1] - E[Y^0|D = 1]) - (E[Y^1|D = 0] - E[Y^0|D = 0])$$

$$= (E[Y^1|D = 1] - E[Y^1|D = 0]) - (E[Y^0|D = 1] - E[Y^0|D = 0])$$

$$= (0) - (0) = 0$$

Randomization

What does randomization do to HTEB?

$$ATT - ATU$$

$$= (E[Y^1|D = 1] - E[Y^0|D = 1]) - (E[Y^1|D = 0] - E[Y^0|D = 0])$$

$$= (E[Y^1|D = 1] - E[Y^1|D = 0]) - (E[Y^0|D = 1] - E[Y^0|D = 0])$$

$$= (0) - (0) = 0$$

HTEB goes to zero!

Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

All we need to estimate the average treatment effect of a policy is:

Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

All we need to estimate the average treatment effect of a policy is:

1. Data on treatment assignment

Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

All we need to estimate the average treatment effect of a policy is:

1. Data on treatment assignment
2. Data on observable outcomes

Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

All we need to estimate the average treatment effect of a policy is:

1. Data on treatment assignment
2. Data on observable outcomes
3. The independence assumption: $Y^1, Y^0 \perp D$

Hands on: randomization and SDO

Let's see how this works in practice by re-constructing our dataset and then randomizing treatment

Hands on: randomization and SDO

Let's see how this works in practice by re-constructing our dataset and then randomizing treatment

```
set.seed(12345)
ate <- 4 # average treatment effect
n_obs <- 100 # number of observations
cons_rand_df <- tibble(
  state = seq(1, n_obs), # state identifier
  Y0 = floor(runif(n_obs)*10)) %>% # control/untreated potential outcome
  mutate(
    D = as.numeric(runif(n()) > 0.5), # randomized treatment
    Y1 = Y0 + ate + round(rnorm(n()))), # generate treatment potential outcome
    Y = D*Y1 + (1-D)*Y0 # generate observed outcome
  ) %>%
  select(
    state, D, Y # keep only observable variables
  )
```

Hands on: randomization and SDO

```
cons_rand_df # data frame of randomized treatment, observable outcome
```

```
## # A tibble: 100 x 3
##   state      D      Y
##   <int> <dbl> <dbl>
## 1     1      0      7
## 2     2      1     11
## 3     3      1     11
## 4     4      1     11
## 5     5      1      8
## 6     6      1      4
## 7     7      1      7
## 8     8      1     11
## 9     9      0      7
## 10    10     0      9
## # ... with 90 more rows
```

Hands on: randomization and SDO

Now take the SDO:

```
cons_rand_df %>%  
  dplyr::group_by(D) %>%  
  dplyr::summarise(meanY = mean(Y))
```

```
## # A tibble: 2 x 2  
##       D   meanY  
##   <dbl> <dbl>  
## 1     0    4.42  
## 2     1    8.82
```

The SDO is $8.82 - 4.42 = 4.40$, a very close estimate of the ATE!

Hands on: randomization and SDO

Now take the SDO:

```
cons_rand_df %>%  
  dplyr::group_by(D) %>%  
  dplyr::summarise(meanY = mean(Y))
```

```
## # A tibble: 2 x 2  
##       D   meanY  
##   <dbl> <dbl>  
## 1     0    4.42  
## 2     1    8.82
```

The SDO is $8.82 - 4.42 = 4.40$, a very close estimate of the ATE!

As $n_{\text{obs}} \rightarrow \infty$, we will have that $SDO \rightarrow ATE$

Randomization: what it does and doesn't do

What does independence imply?

Randomization: what it does and doesn't do

What does independence imply?

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

The two groups have the same potential outcomes, *on average*

Randomization: what it does and doesn't do

What does independence imply?

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

The two groups have the same potential outcomes, *on average*

What does independence **not** imply? That:

$$E[Y^1|D = 1] - E[Y^0|D = 0] = 0 \quad E[Y^1|D = 1] - E[Y^0|D = 1] = 0$$

It does not imply that the observed outcomes are the same across the two groups, nor does it imply that the two potential outcomes of a single group are the same

Independence

Is independence is a reasonable assumption in **observational data**?

Independence

Is independence is a reasonable assumption in **observational data**?

Observational data are just data that we collect from the real world

- e.g. observed state policy choices and outcomes, observed pollution levels and outcomes

Independence

Is independence is a reasonable assumption in **observational data**?

Observational data are just data that we collect from the real world

- e.g. observed state policy choices and outcomes, observed pollution levels and outcomes

Independence is **unlikely** to hold in these scenarios

Independence

Is independence is a reasonable assumption in **observational data**?

Observational data are just data that we collect from the real world

- e.g. observed state policy choices and outcomes, observed pollution levels and outcomes

Independence is **unlikely** to hold in these scenarios

So what can we do?

Conditional independence

Often we will instead rely on **conditional independence**:

$$(Y^1, Y^0 \perp D) | X$$

Conditional independence

Often we will instead rely on **conditional independence**:

$$(Y^1, Y^0 \perp D) | X$$

After controlling for some other variables X , the potential outcomes are independent of treatment $(Y^1, Y^0 \perp D)$

Conditional independence

Often we will instead rely on **conditional independence**:

$$(Y^1, Y^0 \perp D) | X$$

After controlling for some other variables X , the potential outcomes are independent of treatment $(Y^1, Y^0 \perp D)$

This is much weaker than (unconditional) independence: we only need independence to hold for units that share the same X values, not for all units

Conditional independence

Often we will instead rely on **conditional independence**:

$$(Y^1, Y^0 \perp D) | X$$

After controlling for some other variables X , the potential outcomes are independent of treatment $(Y^1, Y^0 \perp D)$

This is much weaker than (unconditional) independence: we only need independence to hold for units that share the same X values, not for all units

Many of the estimation tools used in economics rely on (variants of) the conditional independence assumption